

الجمهورية الشعبية الديمقراطية الجزائرية
People's Democratic Republic of Algeria
وزارة التعليم العالي و البحث العلمي
Ministry of Higher Education and Scientific Research
المدرسة العليا للإعلام الآلي 8 ماي 1945 - سيدي بلعباس
Higher School of Computer Science
8 Mai 1945 - Sidi Bel Abbas



Master's Thesis

To obtain the diploma of Master's Degree

Field of Study: **Computer Science**

Specialization: **Intelligence Artificielle et Science de Données (IASD)**

Theme

A Comprehensive Study of Deep Learning Approaches for Speech-to-Text Recognition

Presented by
BELMANA Soufyane

Defended on: **28 June, 2026**
In front of the jury composed of:

Dr. BEZZAOUCHA Fatima Souad
Dr. CHIKH Asma
Pr. AMAR BENSABER Djamel
Dr. ABDELHAK Soumia

President
Examiner
Supervisor
Incubator Representative

Academic Year: 2025/2026

Acknowledgement

In the name of Allah, the Most Gracious, the Most Merciful.

First, I thank Allah for giving me the strength, patience, and determination to complete this work.

I would like to express my respect and gratitude to my supervisor, Djamel Amar Bensaber, for his supervision and guidance throughout this project.

I also thank the teachers and staff of Higher School of Computer Science of Sidi Bel Abbès for the knowledge and education they provided during my studies.

Finally, I am deeply grateful to my family and friends for their continuous support, encouragement, and prayers during this academic journey.

Abstract

Automatic Speech Recognition (ASR) has undergone a remarkable transformation over the past decade, driven primarily by advances in deep learning. This thesis presents a comprehensive study of deep learning approaches for speech-to-text recognition, tracing the evolution from traditional hybrid systems based on Hidden Markov Models (HMMs) and Gaussian Mixture Models (GMMs) to modern end-to-end architectures powered by Transformers and self-supervised learning.

We systematically analyze the major deep learning architectures that have been applied to ASR, including Deep Neural Networks (DNNs), Convolutional Neural Networks (CNNs), Recurrent Neural Networks (RNNs) with Long Short-Term Memory (LSTM) and Gated Recurrent Unit (GRU) variants, attention mechanisms, and Transformer-based models such as the Conformer. We examine the three main end-to-end paradigms—Connectionist Temporal Classification (CTC), attention-based encoder-decoder models, and RNN-Transducers—along with hybrid approaches that combine their strengths.

Furthermore, we investigate state-of-the-art systems including OpenAI’s Whisper, Meta’s wav2vec 2.0, and Google’s Conformer, analyzing their architectures, training methodologies, and performance characteristics. Our comparative analysis across standard benchmarks demonstrates that modern systems have achieved Word Error Rates (WER) as low as 1.5% on LibriSpeech, surpassing estimated human-level performance. We conclude by discussing remaining challenges including robustness in adverse conditions, multilingual coverage, computational efficiency, and fairness, as well as promising future research directions.

Keywords— Speech Recognition, Deep Learning, Automatic Speech Recognition (ASR), Transformer, Convolutional Neural Networks (CNN), Recurrent Neural Networks (RNN), LSTM, Attention Mechanism, Connectionist Temporal Classification (CTC), wav2vec 2.0, Whisper, Conformer, Self-Supervised Learning, End-to-End Models

Résumé

La reconnaissance automatique de la parole (RAP) a connu une transformation remarquable au cours de la dernière décennie, principalement grâce aux avancées de l'apprentissage profond. Cette thèse présente une étude approfondie des approches d'apprentissage profond pour la reconnaissance vocale (conversion de la parole en texte), retraçant l'évolution depuis les systèmes hybrides traditionnels basés sur les modèles de Markov cachés (HMM) et les modèles de mélange gaussien (GMM) jusqu'aux architectures modernes de bout en bout alimentées par les Transformers et l'apprentissage auto-supervisé.

Nous analysons systématiquement les principales architectures d'apprentissage profond appliquées à la RAP, notamment les réseaux de neurones profonds (DNN), les réseaux de neurones convolutifs (CNN), les réseaux de neurones récurrents (RNN) avec leurs variantes LSTM (Long Short-Term Memory) et GRU (Gated Recurrent Unit), les mécanismes d'attention, ainsi que les modèles basés sur les Transformers tels que le Conformer. Nous examinons les trois principaux paradigmes de bout en bout : la Classification Temporelle Connexionniste (CTC), les modèles encodeur-décodeur à base d'attention, et les RNN-Transducers, ainsi que les approches hybrides combinant leurs forces.

En outre, nous étudions les systèmes à l'état de l'art, notamment Whisper d'OpenAI, wav2vec 2.0 de Meta et le Conformer de Google, en analysant leurs architectures, leurs méthodologies d'entraînement et leurs caractéristiques de performance. Notre analyse comparative sur des benchmarks standards démontre que les systèmes modernes ont atteint des taux d'erreur de mots (WER) aussi bas que 1,5% sur LibriSpeech, dépassant les performances estimées au niveau humain. Nous concluons en discutant les défis restants, notamment la robustesse dans des conditions défavorables, la couverture multilingue, l'efficacité computationnelle et l'équité, ainsi que les directions de recherche futures prometteuses.

Mots-clés : Reconnaissance vocale, Apprentissage profond, Reconnaissance automatique de la parole (RAP), Transformer, Réseaux de neurones convolutifs (CNN), Réseaux de neurones récurrents (RNN), LSTM, Mécanisme d'attention, Classification Temporelle Connexionniste (CTC), wav2vec 2.0, Whisper, Conformer, Apprentissage auto-supervisé, Modèles de bout en bout

ملخص

شهد التعرف التلقائي على الكلام تحولاً ملحوظاً خلال العقد الأخير، مدفوعاً بشكل رئيسي بالتقدم في مجال التعلم العميق. تقدم هذه الأطروحة دراسة شاملة لمقاربات التعلم العميق في التعرف على الكلام وتحويله إلى نص، متتبعاً التطور من الأنظمة الهجينة التقليدية المبنية على نماذج ماركوف المخفية ونماذج المزيج الغاوسي إلى البنى الحديثة من طرف إلى طرف المدعومة بالمحولات والتعلم الذاتي الإشراف. نقوم بتحليل منهجي لأهم بنى التعلم العميق المطبقة في التعرف على الكلام، بما في ذلك الشبكات العصبية العميقة والشبكات العصبية التلافيفية والشبكات العصبية المتكررة بمتغيراتها مثل الذاكرة الطويلة قصيرة المدى ووحدة البوابة المتكررة، وآليات الانتباه، والنماذج المبنية على المحولات مثل الكونفورمر. ندرس ثلاثة نماذج رئيسية من طرف إلى طرف: التصنيف الزمني الاتصالي، ونماذج المشفر-فاك الشفرة المبنية على الانتباه، ومحولات الشبكات العصبية المتكررة، بالإضافة إلى المقاربات الهجينة التي تجمع بين نقاط قوتها. علاوة على ذلك، نستقصي أحدث الأنظمة بما في ذلك ويسبر من أوبن أي آي، وواف تو فيك 0.2 من ميتا، والكونفورمر من غوغل، محللين بنياتها ومنهجيات تدريبها وخصائص أدائها. يُظهر تحليلنا المقارن عبر المعايير القياسية أن الأنظمة الحديثة حققت معدلات خطأ في الكلمات منخفضة تصل إلى 1.5% على مجموعة ليبري سبيتش، متجاوزة الأداء البشري المقدر. نختتم بمناقشة التحديات المتبقية بما في ذلك المتانة في الظروف الصعبة، والتغطية متعددة اللغات، والكفاءة الحسابية، والعدالة، فضلاً عن اتجاهات البحث المستقبلية الواعدة.

الكلمات المفتاحية: التعرف على الكلام، التعلم العميق، التعرف التلقائي على الكلام، المحولات، الشبكات العصبية التلافيفية، الشبكات العصبية المتكررة، الذاكرة الطويلة قصيرة المدى، آلية الانتباه، التصنيف الزمني الاتصالي، Conformer، Whisper، wav2vec 2.0، التعلم الذاتي الإشراف، نماذج من طرف إلى طرف.

	CONTENTS
--	----------

LIST OF FIGURES

LIST OF TABLES

Part I

Introduction

1.1 Context and Motivation

Speech is the most natural and intuitive form of human communication. Since the advent of computing, researchers have sought to bridge the gap between human speech and machine understanding, giving rise to the field of Automatic Speech Recognition (ASR). The ability to convert spoken language into written text has profound implications across numerous domains, including healthcare, education, accessibility, customer service, and human-computer interaction [1].

Traditional ASR systems relied on a complex pipeline of hand-engineered components, including acoustic models based on Hidden Markov Models (HMMs), pronunciation dictionaries, and statistical language models [2]. While these systems achieved reasonable performance in controlled environments, they struggled with variability in accents, background noise, speaker diversity, and spontaneous speech. The development and maintenance of such systems required extensive domain expertise and linguistic resources, limiting their scalability and adaptability.

The emergence of deep learning has fundamentally transformed the landscape of speech recognition. Beginning with the pioneering work of Hinton et al. [3], who demonstrated that Deep Neural Networks (DNNs) could significantly outperform traditional Gaussian Mixture Models (GMMs) in acoustic modeling, the field has witnessed rapid and continuous progress. Subsequent advances in Recurrent Neural Networks (RNNs) [4], Long Short-Term Memory (LSTM) networks [5], Convolutional Neural Networks (CNNs) [6], and most recently, Transformer-based architectures [7] have pushed the boundaries of what is achievable in speech-to-text conversion.

Today, state-of-the-art systems such as OpenAI’s Whisper [8], Meta’s wav2vec 2.0 [9], and Google’s Conformer [10] have achieved near-human-level performance on standard benchmarks, and are being deployed in real-world applications serving billions of users daily.

1.2 Problem Statement

Despite the remarkable progress in deep learning-based speech recognition, several critical challenges remain:

1. **Architectural Diversity:** The rapid proliferation of deep learning architectures—from DNNs and RNNs to Transformers and hybrid models—makes it challenging for researchers and practitioners to understand the relative strengths, weaknesses, and appropriate use cases of each approach.
2. **Robustness:** Many systems still struggle with real-world conditions including background noise, overlapping speakers, accented speech, and domain-specific terminology.

3. **Resource Requirements:** State-of-the-art models often require enormous computational resources for training and may not be suitable for deployment on resource-constrained devices.
4. **Multilingual and Low-Resource Languages:** While significant progress has been made for high-resource languages like English and Mandarin, many of the world’s languages lack sufficient labeled data for training effective models.
5. **End-to-End vs. Hybrid Approaches:** The transition from traditional hybrid systems to fully end-to-end models raises questions about the trade-offs between simplicity, performance, and interpretability.

1.3 Research Objectives

This thesis aims to provide a comprehensive and systematic study of deep learning approaches for speech-to-text recognition. The specific objectives are:

- To present a thorough review of the evolution of speech recognition systems, from traditional approaches to modern deep learning methods.
- To analyze and compare the major deep learning architectures used in ASR, including DNNs, CNNs, RNNs/LSTMs, attention mechanisms, and Transformer-based models.
- To examine end-to-end speech recognition paradigms, including Connectionist Temporal Classification (CTC) [11], attention-based encoder-decoder models, and hybrid CTC-attention approaches.
- To investigate recent advances in self-supervised learning, transfer learning, and multilingual speech recognition.
- To provide a comparative analysis of state-of-the-art systems and their performance on standard benchmarks.
- To discuss open challenges and future research directions in the field.

1.4 Thesis Organization

This thesis is organized into the following parts and chapters:

- **Part I — Introduction:** This current chapter (??) provides the context, motivation, problem statement, and objectives of the thesis.
- **Part II — Background and Fundamentals:** Chapter ?? covers the fundamentals of speech processing and feature extraction. Chapter ?? provides the deep learning foundations necessary for understanding the architectures discussed later.
- **Part III — State of the Art:** Chapter ?? presents a comprehensive review of the state of the art in speech recognition, surveying the evolution from traditional systems to modern deep learning approaches, and comparing leading systems and frameworks.
- **Part IV — Deep Learning Architectures for ASR:** Chapter ?? presents a detailed analysis of the major deep learning architectures applied to speech recognition, including DNNs, CNNs, RNNs/LSTMs, and Transformer-based models.
- **Part V — End-to-End and Modern Approaches:** Chapter ?? examines end-to-end speech recognition paradigms and state-of-the-art systems including Whisper, wav2vec 2.0, and Conformer.

- **Part VI — Comparative Analysis and Discussion:** Chapter ?? provides a comprehensive comparative analysis of the discussed approaches, examining their performance, computational requirements, and practical considerations.
- **Part VII — Conclusion:** Chapter ?? summarizes the key findings and discusses future research directions.

Part II

Background and Fundamentals

CHAPTER 2

SPEECH PROCESSING AND FEATURE EXTRACTION

2.1 Introduction

Before delving into the deep learning architectures used for speech recognition, it is essential to understand the fundamental concepts of speech production, the nature of speech signals, and the feature extraction techniques that transform raw audio into representations suitable for machine learning models. This chapter provides a comprehensive overview of these foundational topics.

2.2 The Nature of Speech Signals

2.2.1 Speech Production

Human speech is produced through a complex coordination of the respiratory system, larynx, and vocal tract. Air expelled from the lungs passes through the vocal folds in the larynx, which may vibrate to produce voiced sounds (such as vowels) or remain open for unvoiced sounds (such as fricatives). The resulting sound wave is then shaped by the resonant properties of the vocal tract—the pharynx, oral cavity, and nasal cavity—which act as acoustic filters [1].

The vocal tract's configuration determines the spectral characteristics of the produced sound, creating distinct patterns for different phonemes. These patterns, known as formants, are the resonant frequencies of the vocal tract and play a crucial role in speech perception and recognition.

2.2.2 Digital Representation of Speech

Speech signals are continuous analog waveforms that must be digitized for computational processing. This digitization involves two key steps:

- **Sampling:** The continuous signal is measured at discrete time intervals. The sampling rate (typically 16 kHz or 44.1 kHz for speech) determines the highest frequency that can be represented, according to the Nyquist theorem.
- **Quantization:** Each sample is represented as a discrete numerical value, typically using 16-bit or 32-bit precision.

The resulting digital signal $x[n]$ represents the amplitude of the speech waveform at each discrete time step n .

Figure ?? illustrates the digitization of a speech signal: (a) continuous analog waveform and (b) discrete sampled representation.




Figure 2.1: Digitization of a speech signal: (a) continuous analog waveform and (b) discrete sampled representation.

2.3 Speech Feature Extraction

Raw audio waveforms contain substantial redundancy and are high-dimensional, making them challenging to process directly. Feature extraction transforms the raw signal into a compact, informative representation that captures the essential characteristics of speech while reducing dimensionality.

2.3.1 Short-Time Fourier Transform (STFT)

Speech signals are non-stationary, meaning their statistical properties change over time. However, over short time intervals (typically 20–25 ms), speech can be considered approximately stationary. The Short-Time Fourier Transform exploits this property by analyzing the signal in overlapping frames:

$$X(m, k) = \sum_{n=0}^{N-1} x[n] \cdot w[n - mH] \cdot e^{-j2\pi kn/N} \quad (2.1)$$

where $w[n]$ is a window function (typically Hamming or Hanning), m is the frame index, H is the hop size, and N is the FFT size. The result is a time-frequency representation known as a spectrogram.

2.3.2 Mel-Frequency Cepstral Coefficients (MFCCs)

MFCCs are among the most widely used features in speech recognition. They are designed to approximate the non-linear frequency perception of the human auditory system using the Mel scale:

$$f_{mel} = 2595 \cdot \log_{10} \left(1 + \frac{f}{700} \right) \quad (2.2)$$

The MFCC extraction pipeline consists of the following steps:

1. **Pre-emphasis:** Apply a high-pass filter to boost high-frequency components: $y[n] = x[n] - \alpha x[n - 1]$, where $\alpha \approx 0.97$.
2. **Framing and Windowing:** Divide the signal into overlapping frames and apply a window function.
3. **FFT:** Compute the power spectrum of each frame.

4. **Mel Filterbank:** Apply a set of triangular bandpass filters spaced according to the Mel scale.
5. **Logarithm:** Take the logarithm of the filterbank energies.
6. **DCT:** Apply the Discrete Cosine Transform to obtain the cepstral coefficients.

Figure ?? illustrates the feature extraction pipeline: from raw audio waveform to spectrogram to MFCC features. The process involves Short-Time Fourier Transform, Mel filterbank application, and Discrete Cosine Transform.




Figure 2.2: Feature extraction pipeline: from raw audio waveform to spectrogram to MFCC features. The process involves Short-Time Fourier Transform, Mel filterbank application, and Discrete Cosine Transform.

Figure ?? illustrates the MFCC feature extraction pipeline showing the sequence of processing steps from raw speech signal to cepstral coefficients.




Figure 2.3: MFCC feature extraction pipeline showing the sequence of processing steps from raw speech signal to cepstral coefficients.

Typically, 13 MFCCs are computed per frame, along with their first and second derivatives (delta and delta-delta coefficients), resulting in a 39-dimensional feature vector per frame.

2.3.3 Mel-Frequency Filterbank Features

In recent years, with the advent of deep learning, log Mel-filterbank features (also called Fbank features) have become increasingly popular as alternatives to MFCCs. These features skip the DCT step of the MFCC pipeline, preserving the full filterbank energy representation. Deep learning models, particularly CNNs and Transformers, can learn the necessary decorrelation and feature transformations directly from data, making the DCT step unnecessary [3].

A typical configuration uses 80 Mel-filterbank channels, producing an 80-dimensional feature vector per frame.

2.3.4 Raw Waveform Features

The most recent trend in speech recognition is to process raw audio waveforms directly, allowing the model to learn optimal feature representations from data. Systems like wav2vec 2.0 [9] and Whisper [8] demonstrate that end-to-end learning from raw or minimally processed audio can achieve state-of-the-art results. This approach eliminates the need for hand-crafted feature engineering and allows the model to discover task-specific representations.

Figure ?? illustrates the evolution of speech feature representations from hand-engineered MFCCs to learned raw waveform features.




Figure 2.4: Evolution of speech feature representations from hand-engineered MFCCs to learned raw waveform features.

2.4 Data Augmentation for Speech

Data augmentation is a critical technique for improving the robustness and generalization of speech recognition systems. Several augmentation strategies have been developed:

2.4.1 Signal-Level Augmentation

- **Noise Addition:** Adding background noise at various signal-to-noise ratios (SNRs) to simulate real-world conditions.
- **Speed Perturbation:** Changing the playback speed (typically $\pm 10\%$) to create variations in speaking rate.
- **Volume Perturbation:** Randomly adjusting the amplitude of the signal.

- **Room Impulse Response (RIR) Simulation:** Convolve the signal with simulated room acoustics to create reverberant speech.

2.4.2 Feature-Level Augmentation: SpecAugment

SpecAugment [12] is a highly effective data augmentation technique that operates on the spectrogram representation. It applies three types of augmentation:

1. **Time Warping:** Applying a random time warping to the spectrogram.
2. **Frequency Masking:** Randomly masking consecutive frequency channels.
3. **Time Masking:** Randomly masking consecutive time steps.

Figure ?? illustrates the SpecAugment data augmentation: (a) original spectrogram, (b) with frequency masking applied, and (c) with time masking applied.




Figure 2.5: SpecAugment data augmentation: (a) original spectrogram, (b) with frequency masking applied, and (c) with time masking applied.

2.5 Standard Datasets and Benchmarks

The development and evaluation of speech recognition systems relies on standardized datasets. Table ?? summarizes the major benchmarks used in the field.

Table 2.1: Major speech recognition datasets and benchmarks.

Dataset	Description	Hours	Language(s)
LibriSpeech	Read English audiobooks	960	English
Common Voice	Crowdsourced multi-language	19,000+	100+
WSJ	Wall Street Journal read speech	80	English
Switchboard	Telephone conversations	300	English
AISHELL-1	Mandarin read speech	170	Mandarin
TED-LIUM	TED talk transcriptions	450	English
GigaSpeech	Multi-source English	10,000	English
VoxPopuli	European Parliament	400K+	23 EU

2.6 Evaluation Metrics

The primary metric for evaluating speech recognition systems is the **Word Error Rate (WER)**, defined as:

$$WER = \frac{S + D + I}{N} \times 100\% \quad (2.3)$$

where S is the number of substitutions, D is the number of deletions, I is the number of insertions, and N is the total number of words in the reference transcript.

Other commonly used metrics include:

- **Character Error Rate (CER)**: Similar to WER but computed at the character level, useful for languages without clear word boundaries.
- **Real-Time Factor (RTF)**: The ratio of processing time to audio duration, measuring computational efficiency.
- **Sentence Error Rate (SER)**: The percentage of sentences with at least one error.

2.7 Conclusion

This chapter has established the foundational concepts necessary for understanding deep learning-based speech recognition. We covered the nature of speech signals, the key feature extraction techniques from traditional MFCCs to modern learned representations, data augmentation strategies, standard datasets, and evaluation metrics. These fundamentals provide the necessary background for the deep learning architectures discussed in the following chapters.

3.1 Introduction

Deep learning, a subset of machine learning based on artificial neural networks with multiple layers of abstraction, has become the dominant paradigm in speech recognition. This chapter provides the theoretical foundations of the key deep learning architectures that have been applied to ASR, forming the basis for the detailed analysis in subsequent chapters.

3.2 Artificial Neural Networks

3.2.1 The Perceptron and Feed-Forward Networks

The fundamental building block of neural networks is the artificial neuron (perceptron), which computes a weighted sum of its inputs and applies a non-linear activation function:

$$y = \sigma \left(\sum_{i=1}^n w_i x_i + b \right) = \sigma(\mathbf{w}^T \mathbf{x} + b) \quad (3.1)$$

where $\mathbf{x}$ is the input vector, $\mathbf{w}$ are the weights, b is the bias, and σ is the activation function.

Figure ?? illustrates the structure of an artificial neuron: inputs are multiplied by weights, summed with a bias, and passed through a non-linear activation function.




Figure 3.1: Structure of an artificial neuron: inputs are multiplied by weights, summed with a bias, and passed through a non-linear activation function.

3.2.2 Activation Functions

The choice of activation function significantly impacts network behavior and training dynamics. The most commonly used activation functions in modern deep learning include:

- **ReLU (Rectified Linear Unit):** $f(x) = \max(0, x)$. The most widely used activation function due to its computational simplicity and effectiveness in mitigating the vanishing gradient problem.
- **Sigmoid:** $\sigma(x) = \frac{1}{1+e^{-x}}$. Maps inputs to the range $(0, 1)$, used in gating mechanisms.
- **Tanh:** $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$. Maps inputs to $(-1, 1)$, commonly used in RNNs.
- **GELU (Gaussian Error Linear Unit):** $f(x) = x \cdot \Phi(x)$. Used in modern Transformer architectures.
- **Softmax:** $\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$. Used in output layers for classification.

Figure ?? illustrates the common activation functions used in deep learning: ReLU, Sigmoid, Tanh, and GELU.




Figure 3.2: Common activation functions used in deep learning: ReLU, Sigmoid, Tanh, and GELU.

3.2.3 Training with Backpropagation

Neural networks are trained using gradient descent and the backpropagation algorithm. The training process involves:

1. **Forward Pass:** Compute the output of the network given an input.
2. **Loss Computation:** Calculate the difference between the predicted and target outputs using a loss function (e.g., cross-entropy for classification).
3. **Backward Pass:** Compute the gradients of the loss with respect to all parameters using the chain rule.
4. **Parameter Update:** Update the parameters using an optimizer such as SGD, Adam, or AdamW.

The Adam optimizer [13], which combines momentum and adaptive learning rates, has become the standard choice for training speech recognition models:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \quad (3.2)$$

where $\hat{m}_t$ and $\hat{v}_t$ are bias-corrected first and second moment estimates of the gradient.

3.3 Deep Neural Networks (DNNs)

Deep Neural Networks consist of multiple stacked hidden layers between the input and output layers. Each layer applies a linear transformation followed by a non-linear activation:

$$\mathbf{h}^{(l)} = \sigma(\mathbf{W}^{(l)}\mathbf{h}^{(l-1)} + \mathbf{b}^{(l)}) \quad (3.3)$$

where $\mathbf{h}^{(l)}$ is the output of layer l , $\mathbf{W}^{(l)}$ and $\mathbf{b}^{(l)}$ are the weight matrix and bias vector, and σ is the activation function.

Figure ?? illustrates the deep Neural Network (DNN) architecture for speech recognition: acoustic features (MFCCs) are input, and the network produces phoneme probability distributions through multiple hidden layers.




Figure 3.3: Deep Neural Network (DNN) architecture for speech recognition: acoustic features (MFCCs) are input, and the network produces phoneme probability distributions through multiple hidden layers.

DNNs were the first deep learning models to achieve breakthrough results in speech recognition [3, 14]. Key advantages include:

- Ability to learn hierarchical representations of acoustic features.
- Superior modeling of complex, non-linear relationships compared to GMMs.

- Effective use of large amounts of training data.

However, standard DNNs process fixed-size input windows and do not inherently model temporal dependencies, limiting their ability to capture the sequential nature of speech.

3.4 Convolutional Neural Networks (CNNs)

CNNs apply learnable filters (kernels) that slide over the input to detect local patterns. A convolutional layer computes:

$$\mathbf{h}_{i,j} = \sigma \left(\sum_m \sum_n \mathbf{W}_{m,n} \cdot \mathbf{x}_{i+m,j+n} + b \right) \quad (3.4)$$

In speech recognition, CNNs are applied to spectrograms or filterbank features, treating them as 2D images where the axes represent time and frequency [6, 15].

Figure ?? illustrates the CNN processing of a speech spectrogram: convolutional filters extract local time-frequency patterns, producing feature maps that are then pooled for dimensionality reduction.




Figure 3.4: CNN processing of a speech spectrogram: convolutional filters extract local time-frequency patterns, producing feature maps that are then pooled for dimensionality reduction.

Key advantages of CNNs for speech:

- **Local Feature Detection:** Ability to capture local spectro-temporal patterns.
- **Translation Invariance:** Robustness to small shifts in frequency or time.
- **Parameter Sharing:** Efficient parameterization through shared filters.
- **Hierarchical Patterns:** Deeper layers capture increasingly complex patterns.

3.5 Recurrent Neural Networks (RNNs)

RNNs are designed to process sequential data by maintaining a hidden state that captures information from previous time steps:

$$\mathbf{h}_t = \sigma(\mathbf{W}_{hx}\mathbf{x}_t + \mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h) \quad (3.5)$$

$$\mathbf{y}_t = \mathbf{W}_{yh}\mathbf{h}_t + \mathbf{b}_y \quad (3.6)$$

This recurrent structure makes RNNs naturally suited for processing speech, which is inherently sequential.

Figure ?? illustrates the recurrent Neural Network unfolded over time steps for speech recognition. The hidden state h_t carries temporal information from previous inputs, enabling the model to capture sequential dependencies in speech.




Figure 3.5: Recurrent Neural Network unfolded over time steps for speech recognition. The hidden state h_t carries temporal information from previous inputs, enabling the model to capture sequential dependencies in speech.

3.5.1 The Vanishing Gradient Problem

Standard RNNs suffer from the vanishing gradient problem, where gradients diminish exponentially as they are propagated through many time steps during backpropagation through time (BPTT). This makes it difficult for RNNs to learn long-range dependencies in sequences [5].

3.5.2 Long Short-Term Memory (LSTM)

The LSTM architecture [5] addresses the vanishing gradient problem through a gating mechanism that controls the flow of information:

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f) \quad (\text{Forget gate}) \quad (3.7)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i) \quad (\text{Input gate}) \quad (3.8)$$

$$\tilde{\mathbf{C}}_t = \tanh(\mathbf{W}_C[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_C) \quad (\text{Candidate}) \quad (3.9)$$

$$\mathbf{C}_t = \mathbf{f}_t \odot \mathbf{C}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{C}}_t \quad (\text{Cell state}) \quad (3.10)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o) \quad (\text{Output gate}) \quad (3.11)$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{C}_t) \quad (\text{Hidden state}) \quad (3.12)$$

Figure ?? illustrates the internal structure of an LSTM cell showing the forget gate (f_t), input gate (i_t), output gate (o_t), cell state (C_t), and hidden state (h_t) with their associated activation functions.




Figure 3.6: Internal structure of an LSTM cell showing the forget gate (f_t), input gate (i_t), output gate (o_t), cell state (C_t), and hidden state (h_t) with their associated activation functions.

3.5.3 Gated Recurrent Unit (GRU)

The GRU [16] is a simplified variant of the LSTM that combines the forget and input gates into a single update gate and merges the cell state and hidden state:

$$\mathbf{z}_t = \sigma(\mathbf{W}_z[\mathbf{h}_{t-1}, \mathbf{x}_t]) \quad (\text{Update gate}) \quad (3.13)$$

$$\mathbf{r}_t = \sigma(\mathbf{W}_r[\mathbf{h}_{t-1}, \mathbf{x}_t]) \quad (\text{Reset gate}) \quad (3.14)$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W}[\mathbf{r}_t \odot \mathbf{h}_{t-1}, \mathbf{x}_t]) \quad (\text{Candidate}) \quad (3.15)$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t \quad (\text{Hidden state}) \quad (3.16)$$

3.5.4 Bidirectional RNNs

In offline speech recognition, the entire utterance is available for processing. Bidirectional RNNs process the sequence in both forward and backward directions, allowing each output to depend on both past and future context:

$$\mathbf{h}_t = [\vec{\mathbf{h}}_t; \overleftarrow{\mathbf{h}}_t] \quad (3.17)$$

Figure ?? illustrates the bidirectional RNN architecture: the forward layer (blue) processes the sequence left-to-right while the backward layer (red) processes right-to-left, combining both directions for each output.




Figure 3.7: Bidirectional RNN architecture: the forward layer (blue) processes the sequence left-to-right while the backward layer (red) processes right-to-left, combining both directions for each output.

3.6 The Attention Mechanism

The attention mechanism [17] allows models to selectively focus on relevant parts of the input when producing each output element. For speech recognition, this enables the decoder to attend to specific portions of the acoustic input when generating each output token.

3.6.1 Scaled Dot-Product Attention

The scaled dot-product attention [7] computes attention weights as:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V} \quad (3.18)$$

where $\mathbf{Q}$, $\mathbf{K}$, and $\mathbf{V}$ are the query, key, and value matrices, and d_k is the dimension of the keys.

3.6.2 Multi-Head Attention

Multi-head attention applies multiple attention functions in parallel, allowing the model to attend to information from different representation subspaces:

$$\text{MultiHead}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) \mathbf{W}^O \quad (3.19)$$

where each $\text{head}_i = \text{Attention}(\mathbf{QW}_i^Q, \mathbf{KW}_i^K, \mathbf{VW}_i^V)$.

Figure ?? illustrates the Multi-head attention mechanism: multiple attention heads operate in parallel on different projections of the queries, keys, and values, then their outputs are concatenated and linearly projected.




Figure 3.8: Multi-head attention mechanism: multiple attention heads operate in parallel on different projections of the queries, keys, and values, then their outputs are concatenated and linearly projected.

3.7 The Transformer Architecture

The Transformer [7] is built entirely on attention mechanisms, dispensing with recurrence and convolutions. It consists of an encoder and decoder, each composed of stacked layers containing multi-head self-attention and position-wise feed-forward networks.

Figure ?? illustrates the the Transformer architecture for speech recognition (ASR), showing the encoder with multi-head self-attention and feed-forward layers, and the decoder with masked self-attention and cross-attention mechanisms.




Figure 3.9: The Transformer architecture for speech recognition (ASR), showing the encoder with multi-head self-attention and feed-forward layers, and the decoder with masked self-attention and cross-attention mechanisms.

3.7.1 Positional Encoding

Since the Transformer has no inherent notion of sequence order, positional encodings are added to the input embeddings:

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{model}}) \quad (3.20)$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{model}}) \quad (3.21)$$

3.7.2 Layer Normalization and Residual Connections

Each sub-layer in the Transformer employs a residual connection followed by layer normalization:

$$\text{output} = \text{LayerNorm}(\mathbf{x} + \text{SubLayer}(\mathbf{x})) \quad (3.22)$$

3.8 Regularization Techniques

Training deep models for speech recognition requires careful regularization to prevent overfitting:

- **Dropout:** Randomly zeroing activations during training with probability p (typically 0.1–0.3).

- **Label Smoothing:** Softening the target distribution to prevent overconfident predictions.
- **Weight Decay:** Adding an L2 penalty to the loss function.
- **SpecAugment:** Data augmentation applied to spectrograms [12].

3.9 Conclusion

This chapter has presented the deep learning foundations essential for understanding modern speech recognition systems. From basic feed-forward networks to sophisticated attention mechanisms and Transformer architectures, each component plays a crucial role in building effective ASR systems. The following chapter will examine how these architectures are specifically applied and adapted for the speech recognition task.

Part III

State of the Art

CHAPTER 4

STATE OF THE ART IN SPEECH-TO-TEXT RECOGNITION

4.1 Introduction

This chapter presents a comprehensive review of the state of the art in speech-to-text recognition. We survey the most influential systems, frameworks, and research contributions that have shaped the field, organizing them chronologically and by paradigm. This review provides the necessary context for understanding the current landscape and identifying the gaps that motivate ongoing research.

4.2 Historical Overview of Speech Recognition

4.2.1 Early Systems (1950s–1980s)

The history of automatic speech recognition dates back to the 1950s. Key milestones include:

- **Audrey (1952)**: Bell Labs’ system that could recognize isolated digits spoken by a single speaker, using formant frequency analysis.
- **Shoebbox (1962)**: IBM’s system capable of recognizing 16 spoken words.
- **HARPY (1976)**: Carnegie Mellon’s system that could handle 1,011 words using a graph-based search approach.
- **Hidden Markov Models (1980s)**: The adoption of HMMs [2] marked a major paradigm shift, establishing the statistical framework that would dominate the field for three decades.

Figure ?? illustrates the major milestones in the history of automatic speech recognition, from the first digit recognizer (1952) to modern LLM-integrated systems (2023).




Figure 4.1: Major milestones in the history of automatic speech recognition, from the first digit recognizer (1952) to modern LLM-integrated systems (2023).

4.2.2 The GMM-HMM Era (1990s–2011)

For nearly three decades, the dominant approach to speech recognition combined Gaussian Mixture Models (GMMs) for acoustic modeling with Hidden Markov Models (HMMs) for sequence modeling [2]. A typical GMM-HMM system consisted of:

1. **Front-end:** Feature extraction (typically MFCCs or PLPs).
2. **Acoustic Model:** GMM-HMM with context-dependent triphone states.
3. **Language Model:** N-gram statistical language model.
4. **Decoder:** Weighted Finite State Transducer (WFST) based search.
5. **Pronunciation Dictionary:** Mapping from words to phoneme sequences.

Figure ?? illustrates the traditional GMM-HMM speech recognition pipeline with five separate components requiring independent optimization.




Figure 4.2: Traditional GMM-HMM speech recognition pipeline with five separate components requiring independent optimization.

While GMM-HMM systems achieved significant success in commercial applications (e.g., telephone-based IVR systems, early dictation software), they suffered from fundamental limitations:

- GMMs have limited modeling capacity for complex, high-dimensional acoustic distributions.
- The pipeline requires expert knowledge to design, train, and maintain each component.
- Feature representations are hand-engineered and may not be optimal for the task.
- Scaling to new languages or domains requires extensive manual effort.

4.3 The Deep Learning Revolution in ASR

4.3.1 DNN-HMM Hybrid Systems (2012)

The deep learning era in speech recognition began with the landmark paper by Hinton et al. [3], which demonstrated that replacing GMMs with DNNs in the acoustic model yielded dramatic improvements. Dahl et al. [14] further advanced this with context-dependent DNN-HMM systems. For a detailed architectural description and mathematical formulation of the hybrid paradigm, see Section ?? in Chapter ??.

Table ?? shows a comparison of the percentage WERs using DNN-HMMs and GMM-HMMs on five different large vocabulary tasks [3].

Table 4.1: A comparison of the percentage WERs using DNN-HMMs and GMM-HMMs on five different large vocabulary tasks [3].

Task	Hours of Training Data	DNN-HMM	GMM-HMM with Same Data	GMM-HMM with More Data
Switchboard (test set 1)	309	18.5%	27.4%	18.6% (2,000 h)
Switchboard (test set 2)	309	16.1%	23.6%	17.1% (2,000 h)
English Broadcast News	50	17.5%	18.8%	—
Bing Voice Search (sentence error rates)	24	30.4%	36.2%	—
Google Voice Input	5,870	12.3%	—	16.0% (>> 5,870 h)
YouTube	1,400	47.6%	52.3%	—

4.3.2 Sequence-to-Sequence Models (2014–2016)

The introduction of sequence-to-sequence models with attention mechanisms marked the beginning of the end-to-end era:

- **CTC (2006/2014)**: Graves et al. [11] introduced CTC, and its combination with deep bidirectional LSTMs [4] produced the first competitive end-to-end systems.
- **Attention-based models (2015)**: Chorowski et al. [18] applied attention mechanisms to speech recognition, enabling flexible alignment learning.
- **Listen, Attend and Spell (2016)**: Chan et al. [19] proposed LAS, the first attention-based model to achieve competitive results on large-vocabulary speech recognition.
- **DeepSpeech 1 & 2 (2014–2016)**: Baidu’s systems [20, 21] demonstrated that simple, scalable architectures trained on large datasets could match complex pipelines.

4.3.3 The Transformer Revolution (2017–2020)

The Transformer architecture [7], originally designed for machine translation, was rapidly adapted for speech recognition:

- **Speech-Transformer (2018)**: Dong et al. [22] applied Transformers to ASR with convolutional subsampling.
- **Transformer Transducer (2020)**: Zhang et al. [23] combined the Transformer encoder with the RNN-T framework for streaming recognition.
- **Conformer (2020)**: Gulati et al. [10] proposed the Conformer, combining convolutions with self-attention, achieving state-of-the-art results on LibriSpeech (1.9%/3.9% WER on test-clean/test-other).

4.4 State-of-the-Art Systems

4.4.1 OpenAI Whisper

Whisper [8] introduced a paradigm shift by training on 680,000 hours of weakly supervised multilingual data. The complete architectural schematic, including the log-Mel spectrogram front-end and cross-attention mechanics, is analyzed in Section ?? of Chapter ?. Key characteristics:

- Standard Transformer encoder-decoder architecture.
- Multitask training: transcription, translation, language ID, timestamps.
- 99 languages supported with a single model.
- Remarkable zero-shot robustness across domains and conditions.
- Model sizes from 39M (Tiny) to 1.55B (Large) parameters.

Figure ?? illustrates the Whisper’s multitask capabilities: a single model handles transcription, translation, language identification, and timestamp prediction.




Figure 4.3: Whisper’s multitask capabilities: a single model handles transcription, translation, language identification, and timestamp prediction.

4.4.2 wav2vec 2.0 and Self-Supervised Learning

The wav2vec 2.0 framework [9] demonstrated that self-supervised pre-training on unlabeled audio can learn powerful speech representations. Its two-stage training process (pre-training with contrastive loss, then fine-tuning with CTC) dramatically reduces the need for labeled data. A detailed breakdown of its feature encoder, quantizer codebooks, and fine-tuning CTC heads is provided in Section ?? of Chapter ?.

4.4.3 HuBERT

HuBERT [24] uses an iterative masked prediction approach with offline clustering to generate pseudo-labels. It achieves results comparable to or better than wav2vec 2.0, particularly in low-resource scenarios. The iterative k-means clustering process and masked predictive training of HuBERT are detailed in Section ?? of Chapter ?.

4.4.4 Google Conformer and USM

Google’s Conformer [10] architecture, which fuses convolution and self-attention in a single block, has become the backbone of many state-of-the-art systems. The block-level diagram and details of the feed-forward, self-attention, and convolution modules of the Conformer are presented in Section ?? of Chapter ?. The Universal Speech Model (USM) [25] scaled this approach to 12 million hours of audio in 300+ languages.

4.4.5 Meta’s Massively Multilingual Speech (MMS)

Meta’s MMS project [26] extended speech technology coverage to over 1,100 languages through:

- Pre-training on 500,000 hours of unlabeled audio in 1,400+ languages.
- Combining self-supervised learning with language-adaptive fine-tuning.
- Demonstrating that scaling data diversity is as important as scaling data quantity.

4.4.6 AssemblyAI Universal Models

AssemblyAI developed a series of production-oriented ASR models—Universal-1, Universal-2, and Universal-3 Pro—based on the Conformer RNN-T architecture [27, 28].

4.4.6.1 Architecture and Pre-training

The core architecture is a **600M-parameter Conformer RNN-T** [27]. The Conformer encoder [10] captures both local acoustic patterns (via convolution) and long-range dependencies (via self-attention). The RNN-T decoder [29] enables streaming inference and reduces hallucination rates relative to autoregressive attention-based decoders. The encoder is pre-trained using **BEST-RQ** [30], a self-supervised method employing a frozen random-projection quantizer—computationally cheaper than iterative clustering approaches such as wav2vec 2.0 or HuBERT. Pre-training data comprises 12.5 million hours of multilingual audio, with fine-tuning on 300,000 hours of supervised and pseudo-labeled data.

4.4.6.2 Key Technical Contributions

Universal-2 [28, 31] introduced two notable innovations:

- **Repeat Token:** A special `<repeat_token>` in the RNN-T vocabulary to handle consecutive identical tokens, achieving up to 90% relative WER reduction on repeated-digit sequences.
- **Neural Text Formatting:** A two-stage post-processing pipeline replacing rule-based formatting—a multi-objective tagging model (punctuation, truecasing, ITN detection) followed by a seq2seq span conversion model.

Universal-3 Pro [32] introduced **natural language prompting** for ASR, allowing users to specify domain terminology, disfluency handling, and custom event tags without model fine-tuning. It also supports real-time streaming with sub-300ms latency.

Table ?? shows the summary of AssemblyAI Universal model family.

Table 4.2: Summary of AssemblyAI Universal model family.

Model	Params	Pre-train Data	Key Contribution	Streaming
Universal-1	600M	12.5M hours	Conformer RNN-T + BEST-RQ SSL	No
Universal-2	600M	12.5M hours	Repeat token, neural text formatting	No
Universal-3 Pro	–	–	Natural language prompting, sub-300ms	Yes

4.5 Comparison of State-of-the-Art Systems

Table ?? provides a comprehensive comparison of the leading speech recognition systems as of 2026.

Table 4.3: Comprehensive comparison of state-of-the-art speech recognition systems.

System	Year	Params	Languages	Training Paradigm	WER*
DeepSpeech 2	2016	50M	2	Supervised, CTC	5.2
ESPnet Conformer	2020	120M	1+	Supervised, CTC-Att	2.1
Conformer-L	2020	118M	1	Supervised, CTC-Att	1.9
wav2vec 2.0-L	2020	300M	1+	Self-supervised + CTC	1.8
HuBERT-XL	2021	1B	1+	Self-supervised + CTC	1.7
Whisper-L v3	2023	1.55B	99	Weakly supervised	2.0
USM	2023	2B	300+	SSL + Supervised	1.5
MMS	2023	1B	1,100+	SSL + Fine-tuning	3.5
AssemblyAI Uni-2	2024	600M	99	SSL + Supervised, RNN-T	~2.0
AssemblyAI Uni-3	2026	–	99+	Promptable RNN-T	–

*WER (%) on LibriSpeech test-clean. “–” indicates not publicly disclosed.

Figure ?? illustrates the quantitative comparison of leading speech recognition systems across their LibriSpeech test-clean Word Error Rate (WER %).




Figure 4.4: Quantitative comparison of leading speech recognition systems across their LibriSpeech test-clean Word Error Rate (WER %).

4.6 Open-Source Toolkits and Frameworks

The rapid progress in speech recognition has been facilitated by several open-source toolkits:

Table ?? shows the major open-source speech recognition toolkits.

Table 4.4: Major open-source speech recognition toolkits.

Toolkit	Organization	Key Features	Framework
Kaldi	JHU	HMM-based, WFST decoder, extensive recipes	C++
ESPnet [33]	CMU/JHU	E2E models, CTC-Att, Conformer, Transducer	PyTorch
Fairseq	Meta	wav2vec 2.0, HuBERT, MMS	PyTorch
NeMo	NVIDIA	Conformer, FastConformer, CTC/RNN-T	PyTorch
Whisper	OpenAI	Multitask, multilingual	PyTorch
WeNet	Community	Production-ready, streaming Conformer	PyTorch
SpeechBrain	Mila	All-in-one speech toolkit	PyTorch

4.7 Recent Trends and Emerging Directions

4.7.1 Integration with Large Language Models

The most recent trend involves integrating ASR with Large Language Models (LLMs). Systems like AudioPaLM and SALMONN use LLMs as decoders for speech understanding tasks, enabling:

- Richer contextual understanding of transcribed speech.
- Joint speech understanding and reasoning.
- Generative error correction of ASR outputs.

4.7.2 Speech Foundation Models

Following the success of foundation models in NLP and vision, speech foundation models are emerging as universal pre-trained representations that can be fine-tuned for diverse downstream tasks beyond recognition, including speaker verification, emotion recognition, and speech translation.

4.7.3 Efficient and On-Device Models

There is growing emphasis on creating efficient models suitable for on-device deployment:

- **Distilled Whisper:** Smaller, faster variants of Whisper through knowledge distillation.
- **FastConformer:** NVIDIA’s optimized Conformer with faster inference.
- **Streaming models:** Real-time capable architectures for voice assistants and live captioning.

Figure ?? illustrates the the current research landscape in automatic speech recognition, showing the major active research directions.




Figure 4.5: The current research landscape in automatic speech recognition, showing the major active research directions.

4.8 Synthesis and Gaps in the Literature

Based on our review of the state of the art, we identify the following key observations and gaps:

1. **Convergence of Approaches:** The best-performing systems increasingly combine elements from multiple paradigms (CNN + Transformer in Conformer, CTC + Attention in hybrid decoders, SSL + supervised fine-tuning).
2. **Data vs. Architecture:** Both Whisper and wav2vec 2.0 demonstrate that data scale and quality can be as important as architectural innovations. AssemblyAI’s Universal-2 further validates this by achieving state-of-the-art results with a modest 600M parameter model trained on 12.5 million hours of data.
3. **Benchmark Saturation:** Standard benchmarks like LibriSpeech are becoming saturated, with top systems achieving below 2% WER. More challenging and diverse benchmarks are needed. AssemblyAI’s emphasis on “last-mile” metrics (proper nouns, formatting, timestamps) represents a shift toward more holistic evaluation.
4. **Real-World Gap:** There remains a significant gap between benchmark performance and real-world deployment scenarios involving noise, accents, domain-specific vocabulary, and conversational speech.
5. **Promptable and Controllable ASR:** The emergence of promptable models like AssemblyAI’s Universal-3 Pro signals a new paradigm where users can guide transcription behavior through natural language, bridging the gap between generic ASR and domain-specific requirements without fine-tuning.
6. **Language Coverage:** Despite impressive progress in multilingual models, the majority of the world’s languages still lack adequate speech recognition support.
7. **Interpretability:** As models become larger and more complex, understanding their internal representations and failure modes becomes increasingly important but challenging.

4.9 Conclusion

This state-of-the-art review has surveyed the major developments in deep learning-based speech recognition, from the initial DNN-HMM breakthrough to the latest self-supervised and multilingual

foundation models. The field has progressed from complex multi-component pipelines requiring extensive domain expertise to elegant end-to-end systems that can be trained with minimal supervision. However, the journey toward truly universal, robust, and efficient speech recognition is far from complete, providing rich opportunities for future research.

Part IV

Deep Learning Architectures for ASR

CHAPTER 5

DEEP LEARNING ARCHITECTURES FOR SPEECH RECOGNITION

5.1 Introduction

This chapter provides a detailed analysis of how the deep learning architectures introduced in Chapter 3 have been specifically applied to the speech recognition task. We trace the evolution from early DNN-HMM hybrid systems to fully end-to-end models, examining the innovations, advantages, and limitations of each approach.

5.2 DNN-HMM Hybrid Systems

The first major breakthrough of deep learning in speech recognition came with DNN-HMM hybrid systems [3, 14]. In this paradigm, DNNs replace the GMM component of traditional HMM-based systems while retaining the HMM framework for sequence modeling.

5.2.1 Architecture

Figure ?? illustrates the DNN-HMM hybrid architecture: the DNN computes posterior probabilities over HMM states, which are then used by the HMM decoder together with a pronunciation dictionary and language model.




Figure 5.1: DNN-HMM hybrid architecture: the DNN computes posterior probabilities over HMM states, which are then used by the HMM decoder together with a pronunciation dictionary and language model.

The DNN takes a context window of acoustic features (typically 11 frames of 40-dimensional filterbank features) and outputs posterior probabilities $P(s|\mathbf{x})$ over HMM states. These posteriors are converted to scaled likelihoods using Bayes' rule:

$$P(\mathbf{x}|s) \propto \frac{P(s|\mathbf{x})}{P(s)} \quad (5.1)$$

5.2.2 Key Results

Hinton et al. [3] demonstrated that DNNs with 4–8 hidden layers of 1024–2048 units could reduce the word error rate by 20–30% relative to GMM-HMM systems across multiple benchmarks. Context-dependent DNN-HMM systems [14] further improved performance by modeling triphone states.

5.3 CNN-Based Approaches

5.3.1 CNNs for Acoustic Modeling

Abdel-Hamid et al. [6] pioneered the application of CNNs to speech recognition, demonstrating that convolutional layers could effectively capture local spectro-temporal patterns. Sainath et al. [15] further showed that deep CNN architectures with multiple convolutional and pooling layers outperform DNNs for large-vocabulary tasks.

Figure ?? illustrates the CNN architecture for speech recognition: convolutional layers extract local features from the spectrogram, followed by fully connected layers for classification.




Figure 5.2: CNN architecture for speech recognition: convolutional layers extract local features from the spectrogram, followed by fully connected layers for classification.

5.3.2 Time-Delay Neural Networks (TDNNs)

TDNNs are a specialized form of 1D CNNs designed for speech recognition. They apply convolutions along the time axis with increasing dilation factors at higher layers, allowing the network to capture both short-term and long-term temporal patterns efficiently.

Figure ?? illustrates the TDNN with increasing dilation factors: each layer captures increasingly larger temporal contexts while maintaining computational efficiency.




Figure 5.3: TDNN with increasing dilation factors: each layer captures increasingly larger temporal contexts while maintaining computational efficiency.

5.4 RNN-Based Approaches

5.4.1 Deep LSTM Networks

Sak et al. [34] demonstrated that deep LSTM networks with multiple stacked layers could significantly outperform DNNs for speech recognition. The key advantage is the LSTM’s ability to capture long-range temporal dependencies in speech.

Graves et al. [4] showed that deep bidirectional LSTM (BLSTM) networks achieve state-of-the-art results by processing speech in both temporal directions. A typical configuration uses 5–7 bidirectional LSTM layers with 512–1024 units per direction.

Figure ?? illustrates the deep bidirectional LSTM network with three stacked layers. Each layer contains forward (blue) and backward (red) LSTM cells, capturing both past and future context.




Figure 5.4: Deep bidirectional LSTM network with three stacked layers. Each layer contains forward (blue) and backward (red) LSTM cells, capturing both past and future context.

5.4.2 CTC-Based RNN Systems

The combination of deep bidirectional LSTMs with Connectionist Temporal Classification (CTC) [11] enabled the first end-to-end speech recognition systems that could directly map acoustic features to character or word sequences without requiring explicit alignment.

The CTC loss function marginalizes over all possible alignments between the input and output sequences:

$$P(\mathbf{y}|\mathbf{x}) = \sum_{\pi \in \mathcal{B}^{-1}(\mathbf{y})} \prod_{t=1}^T P(\pi_t|\mathbf{x}) \quad (5.2)$$

where $\mathcal{B}^{-1}(\mathbf{y})$ is the set of all paths that map to the output $\mathbf{y}$ after removing blanks and repeated characters.

Figure ?? illustrates the CTC alignment and decoding: the model outputs a character (or blank ϵ) for each time step, and the final transcription is obtained by removing blanks and merging repeated characters.




Figure 5.5: CTC alignment and decoding: the model outputs a character (or blank ϵ) for each time step, and the final transcription is obtained by removing blanks and merging repeated characters.

5.5 DeepSpeech and DeepSpeech 2

Baidu’s DeepSpeech [20] and DeepSpeech 2 [21] systems demonstrated that simple, scalable architectures trained on massive datasets could achieve competitive results.

5.5.1 DeepSpeech Architecture

DeepSpeech uses a 5-layer architecture: three fully connected layers, one bidirectional RNN layer, and a final fully connected layer with CTC loss. The key innovation was demonstrating that simpler architectures with sufficient data and compute could rival more complex systems.

5.5.2 DeepSpeech 2

DeepSpeech 2 [21] extended this approach with deeper architectures combining convolutions and bidirectional GRUs:

Figure ?? illustrates the DeepSpeech 2 architecture: convolutional layers for local feature extraction followed by multiple bidirectional GRU layers and CTC output.




Figure 5.6: DeepSpeech 2 architecture: convolutional layers for local feature extraction followed by multiple bidirectional GRU layers and CTC output.

5.6 Attention-Based Encoder-Decoder Models

5.6.1 Listen, Attend and Spell (LAS)

The Listen, Attend and Spell (LAS) model [19] was among the first successful attention-based encoder-decoder models for speech recognition. It consists of three components:

1. **Listener (Encoder):** A pyramidal bidirectional LSTM that processes the input features and reduces the sequence length.
2. **Attention:** A content-based attention mechanism that aligns encoder outputs with decoder steps.
3. **Speller (Decoder):** An LSTM-based decoder that generates the output transcript one character at a time.

Figure ?? illustrates the listen, Attend and Spell (LAS) architecture: the Listener encodes acoustic features using pyramidal BLSTMs, the Attention module aligns encoder outputs with decoder states, and the Speller generates characters autoregressively.




Figure 5.7: Listen, Attend and Spell (LAS) architecture: the Listener encodes acoustic features using pyramidal BLSTMs, the Attention module aligns encoder outputs with decoder states, and the Speller generates characters autoregressively.

5.6.2 Advantages and Limitations

Attention-based models offer several advantages:

- End-to-end training without requiring alignments.
- Implicit language modeling through the decoder.
- Flexible output vocabulary (characters, subwords, or words).

However, they also have limitations:

- Difficulty with very long utterances due to attention mechanisms.
- Left-to-right generation limits parallelization during inference.
- Streaming (online) recognition is challenging.

5.7 Transformer-Based Speech Recognition

5.7.1 Speech-Transformer

Dong et al. [22] adapted the Transformer architecture for speech recognition, demonstrating competitive performance with RNN-based systems. Key adaptations for speech include:

- **Subsampling:** Using convolutional layers to reduce the input sequence length before the Transformer encoder, since speech sequences are much longer than text.
- **Relative Positional Encoding:** Using relative position representations instead of absolute positional encodings to better capture temporal relationships.
- **Chunk-based Processing:** Dividing long utterances into chunks for streaming applications.

5.7.2 Conformer

The Conformer [10] combines the strengths of CNNs and Transformers in a single architecture. Each Conformer block consists of:

1. Feed-Forward Module (first half)

2. Multi-Head Self-Attention Module
3. Convolution Module
4. Feed-Forward Module (second half)

Figure ?? illustrates the conformer block architecture: combines feed-forward modules, multi-head self-attention for global context, and convolution for local pattern extraction, with residual connections throughout.




Figure 5.8: Conformer block architecture: combines feed-forward modules, multi-head self-attention for global context, and convolution for local pattern extraction, with residual connections throughout.

The Conformer achieved state-of-the-art results on LibriSpeech, with a WER of 1.9% on the test-clean set and 3.9% on test-other without an external language model [10].

5.8 Hybrid CTC-Attention Models

Kim et al. [35] proposed a joint CTC-attention framework that combines the strengths of both approaches through multi-task learning:

$$\mathcal{L} = \lambda \mathcal{L}_{\text{CTC}} + (1 - \lambda) \mathcal{L}_{\text{Attention}} \quad (5.3)$$

where λ is a weighting factor (typically 0.3). During decoding, both CTC and attention scores are combined for more robust hypothesis selection.

Figure ?? illustrates the hybrid CTC-Attention architecture: a shared encoder feeds both a CTC branch and an attention-based decoder, with their scores combined during decoding for improved robustness.




Figure 5.9: Hybrid CTC-Attention architecture: a shared encoder feeds both a CTC branch and an attention-based decoder, with their scores combined during decoding for improved robustness.

5.9 Comparison of Architectures

Table ?? summarizes the key characteristics of the major architectural approaches.

Table 5.1: Comparison of deep learning architectures for speech recognition.

Architecture	Strengths	Weaknesses	Best Use Case
DNN-HMM	Proven, interpretable	Requires alignment, complex pipeline	Legacy systems
CNN	Local feature extraction	Limited global context	Feature front-end
BLSTM	Sequential modeling	Slow training, not parallelizable	Offline recognition
Transformer	Parallelizable, global context	Memory-intensive for long sequences	Large-scale training
Conformer	Best of CNN + Transformer	Higher model complexity	State-of-the-art ASR
CTC-Attention	Robust decoding	Increased training complexity	Production systems

5.10 Conclusion

This chapter has traced the evolution of deep learning architectures for speech recognition, from early DNN-HMM hybrid systems to modern Conformer and hybrid CTC-attention models. Each architectural innovation has addressed specific limitations of its predecessors, progressively improving both accuracy and efficiency. The next chapter examines the latest end-to-end systems that have pushed the boundaries of speech recognition performance.

Part V

End-to-End and Modern Approaches

6.1 Introduction

The most significant paradigm shift in speech recognition has been the transition from hybrid systems requiring multiple separately trained components to fully end-to-end models that directly map audio to text. This chapter examines the major end-to-end paradigms and the state-of-the-art systems that have defined the current landscape of speech recognition.

6.2 End-to-End Speech Recognition Paradigms

End-to-end (E2E) models learn to map input speech directly to output text without requiring intermediate representations such as phoneme sequences or alignments. Three main paradigms have emerged: CTC (as detailed in Section ?? of Chapter ??), Attention-based Encoder-Decoder (discussed in Section ?? of Chapter ??), and the RNN-Transducer (RNN-T) introduced below.

Figure ?? illustrates the three main end-to-end speech recognition paradigms: CTC, Attention-based Encoder-Decoder, and RNN-Transducer, each with distinct characteristics.




Figure 6.1: Three main end-to-end speech recognition paradigms: CTC, Attention-based Encoder-Decoder, and RNN-Transducer, each with distinct characteristics.

6.2.1 RNN-Transducer (RNN-T)

The RNN-Transducer [29] extends CTC by incorporating a prediction network (analogous to a language model) that conditions the output on previously emitted labels:

$$P(y_u|t, \mathbf{x}) = \text{Joint}(\text{Encoder}(\mathbf{x})_t, \text{Predictor}(y_{1:u-1})) \quad (6.1)$$

Figure ?? illustrates the RNN-Transducer architecture: the encoder processes audio, the prediction network models label history, and the joint network combines both to produce output probabilities.




Figure 6.2: RNN-Transducer architecture: the encoder processes audio, the prediction network models label history, and the joint network combines both to produce output probabilities.

RNN-T has become the preferred architecture for streaming/online speech recognition in production systems (e.g., Google’s on-device speech recognition) due to its ability to emit outputs incrementally [36].

6.3 Self-Supervised Learning for Speech

Self-supervised learning has emerged as a powerful paradigm for learning speech representations from large amounts of unlabeled audio data, dramatically reducing the need for labeled transcriptions.

6.3.1 wav2vec and wav2vec 2.0

The wav2vec framework [37, 9] learns speech representations through self-supervised pre-training on raw audio. wav2vec 2.0 [9] introduced a particularly effective approach combining contrastive learning with quantized speech representations. It employs a Transformer encoder (as described in Section ?? of Chapter ??) as its backbone architecture and uses a CTC loss function (detailed in Section ?? of Chapter ??) for downstream fine-tuning.

Figure ?? illustrates the wav2vec 2.0 architecture: a CNN feature encoder processes raw audio, masked representations are fed to a Transformer encoder, and contrastive learning with quantized targets enables self-supervised pre-training. Fine-tuning adds a CTC head for speech recognition.




Figure 6.3: wav2vec 2.0 architecture: a CNN feature encoder processes raw audio, masked representations are fed to a Transformer encoder, and contrastive learning with quantized targets enables self-supervised pre-training. Fine-tuning adds a CTC head for speech recognition.

Key results of wav2vec 2.0:

- Achieved a WER of 1.8% on LibriSpeech test-clean with only 960 hours of labeled data.
- With just 10 minutes of labeled data, achieved 4.8% WER on test-clean.
- Demonstrated that pre-training on 53,000 hours of unlabeled audio significantly improves performance.

6.3.2 HuBERT

HuBERT (Hidden-Unit BERT) [24] takes a different approach to self-supervised learning by predicting cluster assignments of masked audio segments:

1. Cluster acoustic features using k-means to generate pseudo-labels.
2. Train a BERT-like model to predict masked pseudo-labels.
3. Iterate: re-cluster using learned representations and retrain.

HuBERT achieved comparable or superior results to wav2vec 2.0, particularly in low-resource settings.

6.3.3 Comparison of Self-Supervised Models

Table ?? shows the comparison of self-supervised speech models on LibriSpeech benchmark.

Table 6.1: Comparison of self-supervised speech models on LibriSpeech benchmark.

Model	Pre-training Method	Unlabeled Data	WER clean	WER other
wav2vec 2.0 Base	Contrastive + Quantization	960h	2.7	5.2
wav2vec 2.0 Large	Contrastive + Quantization	60Kh	1.8	3.3
HuBERT Base	Masked Prediction	960h	2.6	5.0
HuBERT Large	Masked Prediction	60Kh	1.9	3.3
HuBERT X-Large	Masked Prediction	60Kh	1.7	3.1

6.4 OpenAI Whisper

Whisper [8] represents a paradigm shift in speech recognition by demonstrating that large-scale weak supervision can produce remarkably robust models. Building upon the standard Transformer encoder-decoder framework introduced in Section ?? of Chapter ??, the model is designed to optimize sequence-to-sequence translation and transcription tasks.

6.4.1 Architecture

Whisper uses a standard Transformer encoder-decoder architecture with the following specifications:

- **Input:** 80-channel log Mel spectrogram computed from 30-second audio segments.
- **Encoder:** Transformer encoder with two initial convolutional layers for subsampling.
- **Decoder:** Transformer decoder that autoregressively generates tokens.
- **Multitask:** Supports transcription, translation, language identification, and timestamp prediction through special tokens.

Figure ?? illustrates the whisper architecture: a Transformer encoder-decoder model trained on 680,000 hours of weakly supervised data, supporting multiple tasks through special tokens.




Figure 6.4: Whisper architecture: a Transformer encoder-decoder model trained on 680,000 hours of weakly supervised data, supporting multiple tasks through special tokens.

6.4.2 Training Data and Scale

Whisper was trained on 680,000 hours of multilingual audio-text pairs collected from the internet. The model family ranges from 39M to 1.55B parameters:

Table ?? shows the whisper model family: sizes and performance on LibriSpeech test-clean.

Table 6.2: Whisper model family: sizes and performance on LibriSpeech test-clean.

Model	Parameters	Layers	Width	WER (%)
Tiny	39M	4	384	7.6
Base	74M	6	512	5.0
Small	244M	12	768	3.4
Medium	769M	24	1024	2.9
Large	1.55B	32	1280	2.7
Large-v3	1.55B	32	1280	2.0

6.4.3 Key Innovations of Whisper

1. **Weak Supervision at Scale:** Instead of using carefully curated datasets, Whisper leverages noisy, internet-sourced audio-text pairs, trading annotation quality for data quantity.
2. **Zero-Shot Robustness:** Whisper achieves strong performance across diverse domains, accents, and acoustic conditions without fine-tuning.
3. **Multitask Learning:** A single model handles transcription, translation, language detection, and voice activity detection.
4. **Multilingual:** Supports 99 languages with competitive performance.

6.5 Multilingual and Cross-Lingual Models

6.5.1 XLS-R and XLSR

Cross-lingual speech representation learning [38] extends self-supervised pre-training to multiple languages simultaneously. XLS-R models are pre-trained on 436,000 hours of audio in 128 languages, learning universal speech representations that transfer across languages.

6.5.2 Universal Speech Model (USM)

Google’s Universal Speech Model [25] was trained on 12 million hours of audio covering 300+ languages, achieving state-of-the-art results across multiple language families.

6.5.3 Massively Multilingual Speech (MMS)

Meta’s MMS project [26] scaled speech technology to over 1,100 languages by:

- Pre-training on 500,000 hours of unlabeled audio in 1,400+ languages.
- Fine-tuning on labeled data in 1,100+ languages.
- Achieving reasonable recognition quality even for extremely low-resource languages.

Figure ?? illustrates the timeline of major milestones in deep learning-based speech recognition, from DNN-HMM hybrid systems (2012) to large-scale multilingual models (2023).




Figure 6.5: Timeline of major milestones in deep learning-based speech recognition, from DNN-HMM hybrid systems (2012) to large-scale multilingual models (2023).

6.6 Streaming vs. Offline Recognition

A critical distinction in speech recognition systems is whether they operate in *streaming* (online) or *offline* mode.

Table ?? shows the comparison of streaming and offline speech recognition approaches.

Table 6.3: Comparison of streaming and offline speech recognition approaches.

Aspect	Streaming (Online)	Offline
Latency	Real-time output	Full utterance processing
Context	Limited future context	Full bidirectional context
Accuracy	Typically lower	Typically higher
Architecture	RNN-T, Streaming Conformer	BLSTM, Full Transformer
Use Cases	Voice assistants, live captioning	Transcription services, search

6.7 Conclusion

This chapter has examined the cutting-edge end-to-end systems and paradigms that define modern speech recognition. From self-supervised learning frameworks like wav2vec 2.0 and HuBERT to large-scale weakly supervised models like Whisper, the field has made remarkable progress. The next chapter provides a comprehensive comparative analysis of all the approaches discussed throughout this thesis.

Part VI

Comparative Analysis and Discussion

7.1 Introduction

Having examined the major deep learning architectures in Chapter ?? and modern end-to-end systems in Chapter ??, this chapter provides a comprehensive comparative analysis. We evaluate the approaches across multiple dimensions including recognition accuracy, computational efficiency, data requirements, and practical deployment considerations.

7.2 Performance Comparison on Standard Benchmarks

7.2.1 LibriSpeech Results

LibriSpeech [39] is the most widely used benchmark for English speech recognition. Table ?? summarizes the performance of major systems on this benchmark.

Table 7.1: Word Error Rate (%) comparison on LibriSpeech benchmark.

System	Architecture	Year	test-clean	test-other
DNN-HMM (Kaldi)	DNN + HMM	2015	5.3	13.3
DeepSpeech 2	CNN + BiGRU + CTC	2016	5.2	12.7
LAS	BLSTM + Attention	2016	4.1	12.2
Deep BLSTM	BLSTM + CTC	2018	3.8	10.3
Speech-Transformer	Transformer + CTC	2018	2.6	5.7
ContextNet	CNN + CTC	2020	2.1	4.6
Conformer-L	Conformer + CTC-Att	2020	1.9	3.9
wav2vec 2.0 Large	Self-supervised + CTC	2020	1.8	3.3
HuBERT X-Large	Self-supervised + CTC	2021	1.7	3.1
Whisper Large-v3	Transformer Enc-Dec	2023	2.0	4.2
USM	Conformer + SSL	2023	1.5	2.9
<i>Human Performance</i>	—	—	5.8	—

Figure ?? illustrates the progression of Word Error Rate on LibriSpeech test-clean over time. Modern systems have surpassed human-level performance (5.8% WER, dashed red line).




Figure 7.1: Progression of Word Error Rate on LibriSpeech test-clean over time. Modern systems have surpassed human-level performance (5.8% WER, dashed red line).

7.2.2 Key Observations

1. **Rapid Improvement:** WER on LibriSpeech test-clean has decreased from 5.3% (DNN-HMM, 2015) to 1.5% (USM, 2023), a 72% relative reduction.
2. **Surpassing Human Performance:** Modern systems have surpassed the estimated human WER of 5.8% on this benchmark [39].
3. **Diminishing Returns:** The rate of improvement has slowed as performance approaches optimal levels on clean, read speech.
4. **Robustness Gap:** The performance gap between test-clean and test-other remains significant, indicating that noise and speaker variability are still challenging.

7.3 Computational Analysis

7.3.1 Model Size and Complexity

Figure ?? illustrates the Trade-off between model size (parameters) and Word Error Rate on LibriSpeech test-clean. Larger models generally achieve lower WER but with diminishing returns.




Figure 7.2: Trade-off between model size (parameters) and Word Error Rate on LibriSpeech test-clean. Larger models generally achieve lower WER but with diminishing returns.

7.3.2 Training Cost Comparison

Table ?? shows the approximate training costs and requirements for major speech recognition systems.

Table 7.2: Approximate training costs and requirements for major speech recognition systems.

System	GPUs	Training Time	Training Data	Est. Cost
DeepSpeech 2	16 GPUs	3–5 days	12,000h labeled	\$5K
Conformer-L	32 TPUs	3 days	960h labeled	\$10K
wav2vec 2.0 Large	128 GPUs	2.3 days	60Kh unlabeled	\$50K
HuBERT X-Large	256 GPUs	4 days	60Kh unlabeled	\$100K
Whisper Large	256 GPUs	~30 days	680Kh weak	\$500K+
USM	512 TPUs	weeks	12Mh mixed	\$1M+

7.4 Practical Deployment Considerations

7.4.1 On-Device vs. Cloud Deployment

Figure ?? illustrates the deployment strategies for speech recognition systems: cloud-based, on-device, and hybrid approaches with their respective trade-offs.




Figure 7.3: Deployment strategies for speech recognition systems: cloud-based, on-device, and hybrid approaches with their respective trade-offs.

7.4.2 Model Compression Techniques

For deploying large models on resource-constrained devices, several compression techniques are employed:

- **Knowledge Distillation:** Training a smaller “student” model to mimic a larger “teacher” model.
- **Quantization:** Reducing parameter precision from 32-bit to 8-bit or 4-bit.
- **Pruning:** Removing unnecessary weights or entire neurons.
- **Architecture Search:** Automatically finding efficient architectures using Neural Architecture Search (NAS).

Figure ?? illustrates the model compression pipeline: knowledge distillation, quantization, and pruning transform large models into compact versions suitable for edge deployment.




Figure 7.4: Model compression pipeline: knowledge distillation, quantization, and pruning transform large models into compact versions suitable for edge deployment.

7.5 Strengths and Limitations Analysis

7.5.1 Overall Assessment

Table ?? shows the comprehensive strengths and limitations analysis of major approaches.

Table 7.3: Comprehensive strengths and limitations analysis of major approaches.

Approach	Key Strengths	Key Limitations
CTC-based	Simple, monotonic alignment, good for streaming	Conditional independence assumption, needs LM
Attention E2E	Flexible alignment, implicit LM, strong accuracy	Difficult streaming, attention failures on long audio
RNN-T	Streaming capable, strong accuracy, production-proven	Complex training, slower convergence
Self-supervised	Minimal labeled data needed, strong transfer	Massive compute for pre-training, large model size
Whisper	Robust, multilingual, zero-shot capable	Large model, not streaming, high latency

7.6 Remaining Challenges

Despite the remarkable progress, several critical challenges remain:

1. **Noisy and Far-Field Speech:** Performance degrades significantly in reverberant or noisy environments with distant microphones.
2. **Code-Switching:** Handling speech that switches between languages within an utterance remains difficult.
3. **Accented and Dialectal Speech:** Models often perform poorly on underrepresented accents and dialects.
4. **Domain Adaptation:** Transferring models to new domains (e.g., medical, legal) without extensive labeled data.
5. **Computational Efficiency:** Reducing the computational cost of state-of-the-art models for wider deployment.
6. **Hallucination:** End-to-end models can occasionally generate plausible but incorrect text (hallucinations).
7. **Fairness:** Ensuring equitable performance across demographic groups.

7.7 Conclusion

This comparative analysis reveals that deep learning has fundamentally transformed speech recognition, with modern systems achieving super-human performance on standard benchmarks. However, the gap between benchmark performance and real-world robustness, along with computational constraints, presents ongoing challenges. The choice of approach depends heavily on the specific deployment requirements, available resources, and target use cases.

Part VII

Conclusion

8.1 Summary of Contributions

This thesis has presented a comprehensive study of deep learning approaches for speech-to-text recognition, tracing the evolution from traditional systems to the cutting-edge models that define the current state of the art. The key contributions and findings of this work are summarized below.

8.1.1 Evolution of Speech Recognition

We demonstrated that the field of speech recognition has undergone a fundamental transformation driven by deep learning. The transition from hand-engineered, multi-component pipeline systems (GMM-HMM) to unified, end-to-end deep learning models represents one of the most significant paradigm shifts in the history of artificial intelligence. This evolution can be characterized by three major phases:

1. **Hybrid Phase (2012–2015):** Deep Neural Networks replaced Gaussian Mixture Models within the traditional HMM framework, achieving 20–30% relative WER reductions [3, 14].
2. **End-to-End Phase (2015–2020):** CTC-based models [11], attention-based encoder-decoders [19], and RNN-Transducers [29] eliminated the need for separate components, enabling direct audio-to-text mapping.
3. **Foundation Model Phase (2020–present):** Self-supervised pre-training [9], large-scale weak supervision [8], and massively multilingual models [26] have pushed performance beyond human levels on standard benchmarks while dramatically expanding language coverage.

Figure ?? illustrates the the three major phases in the evolution of deep learning-based speech recognition.




Figure 8.1: The three major phases in the evolution of deep learning-based speech recognition.

8.1.2 Architectural Analysis

Our analysis of deep learning architectures revealed distinct strengths for each approach:

- **CNNs** excel at extracting local spectro-temporal features and are computationally efficient, making them ideal as front-end feature extractors.
- **RNNs/LSTMs** effectively model temporal dependencies but suffer from sequential computation bottlenecks and difficulty with very long sequences.
- **Transformers** offer superior parallelization and global context modeling but require careful adaptation for the long sequences typical in speech.
- **Conformers** successfully combine local (CNN) and global (Transformer) processing, achieving state-of-the-art results on multiple benchmarks [10].

8.1.3 Key Findings

The comparative analysis presented in Chapter 6 yielded several important findings:

1. Modern systems have achieved Word Error Rates as low as 1.5% on LibriSpeech test-clean, surpassing the estimated human error rate of 5.8%.
2. Self-supervised pre-training (wav2vec 2.0, HuBERT) can achieve near-state-of-the-art results with as little as 10 minutes of labeled data, dramatically reducing the annotation burden.
3. Large-scale weak supervision (Whisper) produces remarkably robust models that generalize well across domains, accents, and noise conditions without fine-tuning.
4. The choice between streaming and offline architectures remains a critical design decision with significant implications for accuracy, latency, and deployment.
5. Model compression techniques enable deployment of sophisticated models on resource-constrained devices, bridging the gap between research and production.

8.2 Comparative Trade-off Analysis of Leading Systems

While the benchmark results in Chapter 6 establish quantitative performance rankings, a more nuanced analysis of the trade-offs between Whisper, wav2vec 2.0, and Conformer-based systems is essential for guiding architectural choices in practice. Table ?? presents a multi-dimensional comparison.

Table 8.1: Multi-dimensional comparative analysis of four leading ASR architectures [8, 9, 10, 28, 32].

Dimension	Whisper (Large-v3)	wav2vec 2.0 (Large)	Conformer (Large)	AssemblyAI (Uni-2/3)
Architecture	Transformer encoder- decoder	CNN feature encoder + Transformer encoder + CTC head	Conformer encoder (CNN + Trans- former hybrid) + CTC- Attention decoder	Conformer encoder + RNN-T de- coder + Neural Text Formatting seq2seq
Training Paradigm	Weakly super- vised on 680K hours of noisy internet data	Self- supervised contrastive pre-training on 60K hours + CTC fine- tuning	Supervised pre-training on curated la- beled data	SSL pre- training on 12.5M hours + Supervised training on 300K hours
WER (clean / noisy)	2.0% / 4.2%	1.8% / 3.3%	1.9% / 3.9%	~2.0% / 3.8%
Inference Speed	Slow (au- toregressive decoder, RTF ~0.5–1.0)	Moderate (CTC greedy, RTF ~0.1– 0.3)	Fast (CTC or streaming RNN-T, RTF <0.1)	Very Fast (Stream- ing RNN-T, sub-300ms latency, RTF <0.1)
Resource Usage	High (1.55B params, ~10 GB VRAM)	Moderate (300M params, ~4 GB VRAM)	Moderate– Low (118M params, ~2 GB VRAM)	Moderate– High (660M params, cloud API hosted)
Low-Resource / Dialect	Strong zero- shot multi- lingual (99 languages), but weak on under- represented dialects	Excellent transfer to low-resource settings with minimal la- beled data	Requires la- beled data per language; limited mul- tilingual capability without USM- scale training	Strong zero- shot multi- lingual (99+ languages), robust on proper nouns and dialects

The analysis reveals fundamental trade-offs: Whisper offers unmatched multilingual robustness and zero-shot generalization at the cost of high computational overhead and autoregressive latency. wav2vec 2.0 provides the best accuracy-to-data efficiency ratio through self-supervised pre-training, making it particularly valuable for low-resource scenarios. Conformer-based systems achieve the best latency and streaming capability, making them the preferred choice for production deployments requiring real-time inference. AssemblyAI’s Universal models combine the structural benefits of Conformer encoders with the robust streaming capability of RNN-Transducers, delivering industry-leading accuracy on specialized speech and text formatting tasks via hosted APIs.

8.3 Analysis of Hallucinations in End-to-End ASR

A critical failure mode specific to end-to-end ASR systems—particularly autoregressive encoder-decoder models like Whisper—is the phenomenon of *hallucination*: the generation of plausible but acoustically ungrounded text. This section provides a rigorous analysis of its causes and implications.

8.3.1 Architectural Origins of Hallucination

Hallucination in E2E ASR arises from the architectural coupling of acoustic processing and language modeling within a single model. In autoregressive decoder architectures, the output at each time step y_t is conditioned on both the encoder representations $\mathbf{h}$ and all previously generated tokens $y_{<t}$:

$$P(y_t|y_{<t}, \mathbf{h}) = \text{Decoder}(y_{<t}, \text{CrossAttention}(\mathbf{h})) \quad (8.1)$$

When the acoustic signal is ambiguous—due to background noise, extended silence, or low signal-to-noise ratio—the cross-attention weights over the encoder output $\mathbf{h}$ become diffuse. In this regime, the decoder’s implicit language model dominates, causing it to generate text based on *linguistic probability* rather than *acoustic evidence*. This is analogous to a language model operating in “open generation” mode, unconstrained by the input.

8.3.2 Contributing Factors

Several factors exacerbate hallucination in practice:

1. **Attention Collapse on Silence:** During extended pauses or silence, the encoder produces near-uniform representations. The cross-attention mechanism fails to anchor the decoder to specific acoustic events, allowing the autoregressive language prior to generate fluent but fabricated text.
2. **Training Data Bias:** Models trained on weakly supervised data (e.g., Whisper’s 680K hours of internet-sourced pairs) may have learned spurious correlations between audio characteristics and textual patterns, leading to systematic hallucination on certain input types.
3. **Repetition Loops:** The autoregressive nature of the decoder creates a positive feedback loop: once a hallucinated token is generated, it conditions subsequent predictions, potentially triggering cascading errors or repetitive loops.
4. **Long-Form Audio:** For inputs exceeding the model’s context window (30 seconds for Whisper), chunking strategies can introduce boundary artifacts where the decoder hallucinates to “fill” incomplete segments.

8.3.3 Mitigation Strategies

Current approaches to mitigate hallucination include:

- **CTC-based architectures:** By enforcing a monotonic, conditionally independent output distribution, CTC models are structurally incapable of hallucination. RNN-T provides a middle ground with limited hallucination risk.
- **Hybrid CTC-Attention decoding:** Using CTC scores to constrain the attention decoder’s beam search reduces the space for linguistically plausible but acoustically ungrounded hypotheses.

- **Voice Activity Detection (VAD) pre-filtering:** Removing silence segments before inference prevents the decoder from hallucinating on non-speech audio.
- **Confidence-based filtering:** Post-hoc rejection of low-confidence output segments using the decoder’s token-level probabilities.

8.4 Linguistic and Technical Challenges in Arabic and Algerian Speech Recognition

While the systems surveyed in this thesis have achieved remarkable performance on well-resourced languages—particularly English—their application to Arabic, and especially to Algerian dialectal speech, presents unique and formidable challenges that remain largely unresolved.

8.4.1 Challenges of Modern Standard Arabic (MSA)

Modern Standard Arabic poses several inherent difficulties for ASR systems [40]:

- **Rich Morphology:** Arabic is a morphologically rich language with a root-and-pattern system. A single triliteral root (e.g., *k-t-b*, “to write”) generates dozens of derived forms through prefixation, suffixation, and infixation, dramatically expanding the vocabulary space relative to English.
- **Diacritization Ambiguity:** Short vowels (diacritics) are typically omitted in written Arabic, creating extensive homographic ambiguity. The word *ktb* (ك ت ب) can be read as *kataba* (“he wrote”), *kutub* (“books”), *kutiba* (“it was written”), among others. This ambiguity complicates language modeling for ASR.
- **Phonetic Complexity:** Arabic contains emphatic consonants (/s /, /d /, /t /, /ð /), pharyngeal fricatives (/ħ/, /ʕ /), and a uvular stop (/q/) that are absent in most European languages and are often poorly modeled by multilingual systems.
- **Diglossia:** Arabic-speaking populations use MSA in formal contexts but communicate daily in regional dialects that differ substantially from MSA in phonology, morphology, and lexicon, creating a fundamental mismatch between training data (predominantly MSA) and real-world speech.

8.4.2 Algerian Darja: A Particularly Challenging Dialect

Algerian Arabic (*Darja*) presents additional challenges beyond those of MSA [41]:

- **Lack of Standardized Orthography:** Darja has no official written form. Speakers write it using Arabic script, Latin script (Arabizi), or a mix of both, with no consistent spelling conventions. This makes it extremely difficult to create standardized transcription corpora for ASR training.
- **Phonetic Divergence:** Algerian Darja exhibits significant phonological shifts from MSA, including the merger or shift of classical phonemes (e.g., /q/ → [] or [g] depending on region), vowel reduction, and consonant cluster simplifications influenced by centuries of Berber, French, Turkish, and Spanish contact.
- **Regional Variation:** Algeria’s vast geography produces substantial regional variation. The speech of Oran, Algiers, Constantine, and southern Saharan regions can differ significantly in vocabulary, phonology, and intonation, effectively constituting distinct sub-dialects.

- **Lexical Borrowing:** Darja incorporates extensive loanwords from French (*tabla* for “table,” *tronportir* for “to transport”), Berber, and Turkish, creating a lexicon unrecognizable to MSA-trained language models.

8.4.3 The Code-Switching Challenge

Perhaps the most distinctive challenge for ASR in Algeria is pervasive **code-switching**—the fluid alternation between Arabic, French, and sometimes Berber within a single utterance or even a single sentence [42]. For example, a speaker might say: “*Rahi tkhdem bien, mais lazem nadirou une réunion bach nreglou el problème*” (“It’s working well, but we need to hold a meeting to fix the problem”)—mixing Darja, French nouns, and French verbs within one sentence.

This poses severe challenges for conventional ASR:

- **Language Model Confusion:** Standard monolingual language models cannot predict cross-lingual transitions, leading to high perplexity at switch points.
- **Acoustic Model Mismatch:** Phoneme inventories differ between languages; a model trained primarily on one language will misrecognize phonemes from the other.
- **Tokenization Ambiguity:** Subword tokenizers (BPE, SentencePiece) trained on a single language produce fragmented, meaningless tokens when applied to code-switched speech.

8.4.4 Dataset Scarcity

The fundamental bottleneck for Arabic dialectal ASR is the severe lack of annotated speech data [41]. While English benefits from thousands of hours of curated corpora (LibriSpeech, Common Voice, GigaSpeech), Algerian Darja has virtually no publicly available, high-quality transcribed speech datasets. Existing Arabic corpora (e.g., MGB-2, GALE Arabic) focus on MSA broadcast speech, which differs substantially from spontaneous dialectal conversation. This data poverty prevents effective supervised training and limits the utility of self-supervised approaches like wav2vec 2.0 that still require labeled fine-tuning data.

8.5 Toward an Experimental Methodology for ASR Benchmarking

To complement the theoretical analysis presented in this thesis, we outline a reproducible experimental methodology for benchmarking open-source ASR models on Arabic speech data.

8.5.1 Experimental Framework

The proposed experimental setup evaluates models across three dimensions:

1. **Accuracy:** Measured by Word Error Rate (WER), computed as:

$$\text{WER} = \frac{S + D + I}{N} \times 100\% \quad (8.2)$$

where S , D , I , and N denote substitutions, deletions, insertions, and total reference words, respectively.

2. **Latency:** Measured as inference time per audio second (Real-Time Factor):

$$\text{RTF} = \frac{T_{\text{inference}}}{T_{\text{audio}}} \quad (8.3)$$

An $\text{RTF} < 1$ indicates faster-than-real-time processing.

3. **Resource Utilization:** Peak GPU/CPU memory consumption during inference.

8.5.2 Proposed Setup

Table ?? shows the proposed experimental configuration for ASR benchmarking.

Table 8.2: Proposed experimental configuration for ASR benchmarking.

Component	Specification
Models	Whisper (tiny, base, small), wav2vec 2.0 (base, large), Conformer (via NeMo)
Dataset	Mozilla Common Voice 13.0 (Arabic subset, test split)
Framework	Python 3.10+, Hugging Face <code>transformers</code> , <code>evaluate</code> (WER metric)
Hardware	NVIDIA GPU (e.g., RTX 3060 or above), 16 GB RAM minimum
Metrics	WER (%), RTF (seconds), Peak VRAM (MB)

This methodology enables a controlled, reproducible comparison of state-of-the-art models on Arabic speech, and can be extended to dialectal data as annotated Darja corpora become available.

8.6 Future Research Directions

Based on our comprehensive analysis, we identify the following promising directions for future research:

8.6.1 Multimodal Speech Recognition

Integrating visual information (lip reading, gestures) with audio for audio-visual speech recognition promises improved robustness in noisy environments. Recent advances in vision-language models suggest that multimodal approaches could significantly enhance recognition accuracy in challenging conditions.

8.6.2 Large Language Model Integration

The integration of large language models (LLMs) with speech recognition systems represents a promising frontier. By leveraging the powerful language understanding capabilities of models like GPT-4 and Gemini, speech recognition could benefit from:

- Improved contextual understanding and disambiguation.
- Better handling of rare words and domain-specific terminology.
- Enhanced error correction through linguistic reasoning.

8.6.3 Efficient Architectures

Developing more computationally efficient architectures that maintain high accuracy while reducing resource requirements is crucial for democratizing speech recognition technology. Promising approaches include:

- Linear attention mechanisms that reduce the quadratic complexity of standard Transformers.
- State-space models (e.g., Mamba) as alternatives to attention for sequence modeling.
- Hardware-aware neural architecture search for optimizing model design for specific deployment targets.

8.6.4 Universal Speech Models

Building truly universal speech models that work across all languages, accents, domains, and acoustic conditions remains an ambitious but achievable goal. Recent efforts like Google’s USM [25] and Meta’s MMS [26] have made significant strides, but coverage of the world’s 7,000+ languages is still far from complete.

8.6.5 Responsible AI and Fairness

Ensuring that speech recognition systems perform equitably across all demographic groups—regardless of accent, dialect, age, or gender—is an increasingly important research direction. This includes:

- Developing benchmarks that evaluate fairness across diverse populations.
- Creating methods to detect and mitigate bias in training data and models.
- Establishing guidelines for responsible deployment of speech technology.

8.7 Concluding Remarks

Deep learning has fundamentally transformed speech-to-text recognition from a research problem to a technology that billions of people use daily. The progress over the past decade has been extraordinary: systems that once required carefully controlled acoustic environments and extensive manual engineering now work robustly in diverse real-world conditions, across dozens of languages, and on devices ranging from smartphones to cloud servers.

Yet significant challenges remain. The gap between benchmark performance and real-world robustness, the computational cost of state-of-the-art models, the coverage of low-resource languages, and the imperative for fairness all represent active frontiers of research. As the field continues to evolve—driven by advances in model architectures, training paradigms, and computational infrastructure—we can anticipate speech recognition systems that are even more accurate, efficient, inclusive, and ubiquitous.

This thesis has aimed to provide researchers and practitioners with a comprehensive understanding of the deep learning landscape for speech recognition, enabling informed decisions about architectural choices, training strategies, and deployment considerations. We hope this work serves as both a reference for current approaches and a foundation for future innovations in this dynamic and impactful field.

REFERENCES

- [1] Daniel Jurafsky and James H. Martin. *Speech and Language Processing*. 3rd. Pearson, 2024. URL: https://web.stanford.edu/~jurafsky/slp3/ed3book_jan26.pdf.
- [2] Lawrence R. Rabiner. “A Tutorial on Hidden Markov Models and Selected Applications in Speech Recognition.” In: *Proceedings of the IEEE* 77.2 (1989), pp. 257–286. URL: <https://www.cs.ubc.ca/~murphyk/Bayes/rabiner.pdf>.
- [3] Geoffrey Hinton et al. “Deep Neural Networks for Acoustic Modeling in Speech Recognition.” In: *IEEE Signal Processing Magazine* 29.6 (2012), pp. 82–97. URL: <https://www.cs.toronto.edu/~hinton/absps/DNN-2012-proof.pdf>.
- [4] Alex Graves, Abdel-rahman Mohamed, and Geoffrey Hinton. “Speech Recognition with Deep Recurrent Neural Networks.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2013), pp. 6645–6649. URL: <https://www.cs.toronto.edu/~fritz/absps/RNN13.pdf>.
- [5] Sepp Hochreiter and Jürgen Schmidhuber. “Long Short-Term Memory.” In: *Neural Computation* 9.8 (1997), pp. 1735–1780. URL: https://www.researchgate.net/publication/13853244_Long_Short-Term_Memory.
- [6] Ossama Abdel-Hamid et al. “Convolutional Neural Networks for Speech Recognition.” In: *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 22.10 (2014), pp. 1533–1545. URL: https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/CNN_ASLPTrans2-14.pdf.
- [7] Ashish Vaswani et al. “Attention Is All You Need.” In: *Advances in Neural Information Processing Systems* 30 (2017), pp. 5998–6008. URL: https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf.
- [8] Alec Radford et al. “Robust Speech Recognition via Large-Scale Weak Supervision.” In: *Proceedings of the 40th International Conference on Machine Learning* (2023), pp. 28492–28518. URL: <https://cdn.openai.com/papers/whisper.pdf>.
- [9] Alexei Baevski et al. “wav2vec 2.0: A Framework for Self-Supervised Learning of Speech Representations.” In: *Advances in Neural Information Processing Systems* 33 (2020), pp. 12449–12460. URL: <https://proceedings.neurips.cc/paper/2020/file/92d1e1eb1cd6f9fba3227870bb6d7f07-Paper.pdf>.
- [10] Anmol Gulati et al. “Conformer: Convolution-augmented Transformer for Speech Recognition.” In: *Proceedings of Interspeech* (2020), pp. 5036–5040. URL: https://www.isca-archive.org/interspeech_2020/gulati20_interspeech.pdf.

- [11] Alex Graves et al. “Connectionist Temporal Classification: Labelling Unsegmented Sequence Data with Recurrent Neural Networks.” In: *Proceedings of the 23rd International Conference on Machine Learning* (2006), pp. 369–376. URL: <https://people.idsia.ch/~santiago/papers/icml2006.pdf>.
- [12] Daniel S. Park et al. “SpecAugment: A Simple Data Augmentation Method for Automatic Speech Recognition.” In: *Proceedings of Interspeech* (2019), pp. 2613–2617. URL: https://www.isca-archive.org/interspeech_2019/park19e_interspeech.pdf.
- [13] Diederik P. Kingma and Jimmy Ba. “Adam: A Method for Stochastic Optimization.” In: *Proceedings of the 3rd International Conference on Learning Representations (ICLR)* (2015). URL: <https://www.intel.com/content/dam/www/public/us/en/ai/documents/1412.6980.pdf>.
- [14] George E. Dahl et al. “Context-Dependent Pre-Trained Deep Neural Networks for Large-Vocabulary Speech Recognition.” In: *IEEE Transactions on Audio, Speech, and Language Processing* 20.1 (2012), pp. 30–42. URL: <https://www.cs.toronto.edu/~gdahl/papers/DBN4LVCSR-TransASLP.pdf>.
- [15] Tara N. Sainath et al. “Deep Convolutional Neural Networks for LVCSR.” In: *Neural Networks* 64 (2013), pp. 8614–8618. URL: https://www.cs.toronto.edu/~asamir/papers/icassp13_cnn.pdf.
- [16] Kyunghyun Cho et al. “Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation.” In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (2014), pp. 1724–1734. URL: <https://arxiv.org/pdf/1406.1078>.
- [17] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. “Neural Machine Translation by Jointly Learning to Align and Translate.” In: *Proceedings of the 3rd International Conference on Learning Representations (ICLR)* (2015). URL: <https://arxiv.org/pdf/1409.0473>.
- [18] Jan K. Chorowski et al. “Attention-Based Models for Speech Recognition.” In: *Advances in Neural Information Processing Systems* 28 (2015), pp. 577–585. URL: <https://arxiv.org/pdf/1506.07503>.
- [19] William Chan et al. “Listen, Attend and Spell: A Neural Network that Learns to Spell.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2016). URL: <https://arxiv.org/pdf/1508.01211>.
- [20] Awni Hannun et al. “Deep Speech: Scaling up end-to-end speech recognition.” In: *arXiv preprint arXiv:1412.5567* (2014). URL: <https://arxiv.org/pdf/1412.5567>.
- [21] Dario Amodei et al. “Deep Speech 2: End-to-End Speech Recognition in English and Mandarin.” In: *Proceedings of the 33rd International Conference on Machine Learning* (2016), pp. 173–182. URL: <https://arxiv.org/pdf/1512.02595>.
- [22] Linhao Dong, Shuang Xu, and Bo Xu. “Speech-Transformer: A No-Recurrence Sequence-to-Sequence Model for Speech Recognition.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2018), pp. 5884–5888. URL: https://houwx.net/files/papers/others/2018_icassp_speech_transformer.pdf.
- [23] Qian Zhang et al. “Transformer Transducer: A Streamable Speech Recognition Model with Transformer Encoders and RNN-T Loss.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2020), pp. 7829–7833. URL: <https://arxiv.org/pdf/2002.02562>.

- [24] Wei-Ning Hsu et al. “HuBERT: Self-Supervised Speech Representation Learning by Masked Prediction of Hidden Units.” In: *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 29 (2021), pp. 3451–3460. URL: <https://arxiv.org/pdf/2106.07447>.
- [25] Yu Zhang et al. “Google USM: Scaling Automatic Speech Recognition Beyond 100 Languages.” In: *arXiv preprint arXiv:2303.01037* (2023). URL: <https://arxiv.org/pdf/2303.01037>.
- [26] Vineel Pratap et al. “Scaling Speech Technology to 1,000+ Languages.” In: *arXiv preprint arXiv:2305.13516* (2023). URL: <https://arxiv.org/pdf/2305.13516>.
- [27] Francis McCann Ramirez et al. “Anatomy of Industrial Scale Multilingual ASR.” In: *arXiv preprint arXiv:2404.09841* (2024). AssemblyAI Research. URL: <https://arxiv.org/pdf/2404.09841>.
- [28] AssemblyAI Research. *Universal-2: Solving the Last Mile Issues in Speech-to-Text*. <https://www.assemblyai.com/research/universal-2>. Technical report, AssemblyAI. 2024. URL: <https://www.assemblyai.com/research/universal-2>.
- [29] Alex Graves. “Sequence Transduction with Recurrent Neural Networks.” In: *arXiv preprint arXiv:1211.3711* (2012). URL: <https://arxiv.org/pdf/1211.3711>.
- [30] Chung-Cheng Chiu et al. “Self-Supervised Learning with Random-Projection Quantizer for Speech Recognition.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2022), pp. 3663–3667. URL: <https://arxiv.org/pdf/2202.01855>.
- [31] Yash Khare et al. “Universal-2-TF: Robust All-Neural Text Formatting for ASR.” In: *arXiv preprint* (2025). AssemblyAI Research. URL: <https://arxiv.org/pdf/2501.05948>.
- [32] AssemblyAI Research. *Universal-3 Pro: Promptable Voice AI for Speech Recognition*. <https://www.assemblyai.com/universal-3-pro>. Technical announcement, AssemblyAI. 2026. URL: <https://www.assemblyai.com/universal-3-pro>.
- [33] Shinji Watanabe et al. “ESPnet: End-to-End Speech Processing Toolkit.” In: *Proceedings of Interspeech* (2018), pp. 2207–2211. URL: <https://arxiv.org/pdf/1804.00015>.
- [34] Haşim Sak, Andrew Senior, and Françoise Beaufays. “Long Short-Term Memory Recurrent Neural Network Architectures for Large Scale Acoustic Modeling.” In: *Proceedings of Interspeech* (2014), pp. 338–342. URL: <https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/43905.pdf>.
- [35] Suyoun Kim, Takaaki Hori, and Shinji Watanabe. “Joint CTC-Attention Based End-to-End Speech Recognition Using Multi-task Learning.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2017), pp. 4835–4839. URL: <https://arxiv.org/pdf/1609.06773>.
- [36] Yanzhang He et al. “Streaming End-to-End Speech Recognition for Mobile Devices.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2019), pp. 6381–6385. URL: <https://arxiv.org/pdf/1811.06621>.
- [37] Steffen Schneider et al. “wav2vec: Unsupervised Pre-Training for Speech Recognition.” In: *Proceedings of Interspeech* (2019), pp. 3465–3469. URL: <https://arxiv.org/pdf/1904.05862>.

- [38] Alexis Conneau et al. “Unsupervised Cross-lingual Representation Learning for Speech Recognition.” In: *Proceedings of Interspeech* (2020). URL: <https://arxiv.org/pdf/2006.13979>.
- [39] Vassil Panayotov et al. “Librispeech: An ASR corpus based on public domain audio books.” In: *Proceedings of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (2015), pp. 5206–5210. URL: https://www.danielpovey.com/files/2015_icassp_librispeech.pdf.
- [40] Fadi Biadisy, Julia Hirschberg, and Nizar Habash. “Spoken Arabic Dialect Identification Using Phonotactic Modeling.” In: *Proceedings of the EACL 2009 Workshop on Computational Approaches to Semitic Languages* (2009), pp. 53–61. URL: <https://aclanthology.org/W09-0807.pdf>.
- [41] Mohamed Maamouri et al. “Developing and Using a Pilot Dialectal Arabic Treebank.” In: *Proceedings of the Fifth International Conference on Language Resources and Evaluation (LREC)* (2006). URL: http://www.lrec-conf.org/proceedings/lrec2006/pdf/543_pdf.pdf.
- [42] Djegdjiga Amazouz, Martine Adda-Decker, and Gilles Adda. “Addressing Code-Switching in French/Algerian Arabic Speech.” In: *Proceedings of Interspeech* (2018), pp. 62–66. URL: https://shs.hal.science/halshs-01969148/file/Addressing_Code-Switching_in_FrenchAlgerian_Arabic.pdf.